Natural Language Processing (NLP)

What is NLP?

Natural Language Processing (NLP) is a field at the intersection of linguistics, computer science, and artificial intelligence that enables machines to understand, interpret, and generate human language.

Example:

- Google Translate
- ChatGPT
- Alexa
- spam filters
- sentiment analysis.

NLU (Natural Language Understanding)

Definition:
NLU is the part of NLP that focuses on understanding and interpreting the meaning, structure, and intent behind human language.

Goal: Convert unstructured text into a machine-understandable format.

Key Sub-Tasks of NLU:

Task	Description	Example
Intent Recognition	Detecting the user's purpose	"Book me a flight" → Intent: BookFlight
Entity Recognition	Identifying important keywords or phrases	"Book a flight to *Delhi*" → Entity: Delhi
Named Entity Recognition (NER)	Extracting names of people, places, dates, etc.	"Meet *John* on *Friday*"
Part-of-Speech Tagging	Assigning word types (noun, verb, adj...)	"She runs" → She/PRONOUN, runs/VERB
Dependency Parsing	Understanding grammatical structure and word relations	"The boy kicked the ball" → Subject → Verb
Coreference Resolution	Resolving pronouns or repeated references	"John said he is coming" → "he" = "John"
Sentiment Analysis	Determining emotional tone	"I love this product" → Positive

Task	Description	Example
Text Classification	Assigning labels to a sentence or document	Spam or Not Spam
Semantic Role Labeling	Identifying roles like who did what to whom	"Mary gave John a book" → Mary=Giver

Tools and Models for NLU:

- **Rule-based**: Regex, spaCy patterns
- **ML-based**: SVM, Random Forest, CRF
- **DL-based**: LSTMs, BiLSTM+CRF
- **Transformer-based**: BERT, RoBERTa, DistilBERT

NLG (Natural Language Generation)

Definition:
NLG is the task of generating human-like text from structured or unstructured data.

Goal: Convert machine-readable data into natural language.

Key Sub-Tasks of NLG:

Task	Description	Example
Content Planning	Deciding what information to include	"User bought 3 items worth $60"
Sentence Planning	Organizing how content will be structured grammatically	"You purchased 3 items for $60."
Surface Realization	Generating the actual sentence from plan	"Thank you for your purchase of $60."
Text Summarization	Generating a short version of a long text	TL;DR of an article
Data-to-Text	Describing structured data in plain language	Weather: "23°C, cloudy" → "It is cloudy today."
Question Generation	Creating questions from context	Input: "Apple is a company." → "What is Apple?"
Conversational Responses	Generating replies in chatbots	User: "How's the weather?" → Bot: "It's sunny!"
Story Generation	Long-form creative writing	Tools like GPT, Claude, etc.

Tools and Models for NLG:

- **Templates**: Predefined sentence structures.
- **Statistical**: N-gram models, Markov chains.
- **Deep Learning**: RNNs, LSTM
- **Transformer-based**: GPT-2/3/4, T5, BART

Relationship Between NLU and NLG

Component	Function	Example
NLU	Understand input	User: "Book a flight to Delhi"
NLG	Generate a meaningful response	Bot: "Sure, when would you like to travel?"

In a chatbot:

- **NLU** identifies intent = BookFlight, entity = Delhi
- **NLG** constructs the response dynamically.

Example: Voice Assistant Workflow

User: "Remind me to call mom at 6 PM"

1. NLU:
   → Intent: CreateReminder
   → Entity: "call mom", "6 PM"

2. Dialogue Manager:
   → Fills reminder slot, asks for date if needed

3. NLG:
   → Response: "Okay, I'll remind you to call mom at 6 PM."

Term	Full Form	Focus
NLU	Natural Language Understanding	**Interpreting** human input
NLG	Natural Language Generation	**Generating** human-like text

Together, **NLU + NLG** power:

- Chatbots
- Voice Assistants
- Virtual Agents
- Customer Support AI
- Smart Home Systems

Key Tasks in NLP

Task	Description	Example
Tokenization	Splitting text into words/tokens	"I love NLP" → ["I", "love", "NLP"]
Part-of-Speech (POS) Tagging	Identifying grammatical categories	"I love NLP" → [PRP, VBP, NNP]
Named Entity Recognition (NER)	Recognizing real-world entities	"Barack Obama was born in Hawaii" → ["PERSON", "GPE"]
Text Classification	Classifying documents or messages	Spam vs Not Spam
Sentiment Analysis	Detecting emotion in text	"I hate this!" → Negative
Text Summarization	Shortening a long document	News highlights
Machine Translation	Translating between languages	English → French
Question Answering	Answering based on a document	"Who is president of India?"
Chatbots	Interactive conversations	Support assistants

Text Preprocessing Techniques

Tokenization

Breaks text into individual words or sub-words.

- **Word Tokenization:** “I love NLP” → [“I”, “love”, “NLP”]
- **Sentence Tokenization:** Splits paragraph into sentences.

Lowercasing

Converts text to lowercase to avoid case mismatch.

"NLP" and "nlp" → both become "nlp"

Stopword Removal

Removes common words that carry little meaning (e.g., "the", "is", "in").

Stemming

Reduces words to their root form.

"playing" → "play"

Library: PorterStemmer, SnowballStemmer

Lemmatization

More intelligent root-word finding than stemming.

"better" → "good"

Punctuation & Special Character Removal

"I ❤NLP!!!" → "I NLP"

Text Representation (Vectorization)

Computers need numbers — here's how we convert text into numerical form:

Bag of Words (BoW)

Counts word frequency.

Text	"NLP"	"is"	"fun"
NLP is fun	1	1	1
NLP is powerful	1	1	0

TF-IDF (Term Frequency–Inverse Document Frequency)

Highlights **important words** in a document.

Formula:
TF = (No. of times term appears) / (Total terms)
IDF = log(Total docs / Docs containing term)
TF-IDF = TF × IDF

Word Embeddings

Context-aware vector representations.

- **Word2Vec** (Google)
- **GloVe** (Stanford)
- **FastText** (Facebook)
- **BERT Embeddings** (Transformers)

Example:
"king" - "man" + "woman" ≈ "queen"

Language Models

Language Models (LM) predict the next word in a sequence.

N-Gram Models

Probabilistic model based on n previous words.

Formula (Bigram):
P(w2 | w1) = Count(w1 w2) / Count(w1)

Limitation: Sparse and memory-intensive.

Neural Language Models

Use deep learning.

- **RNN (Recurrent Neural Networks)**
- **LSTM / GRU** – Handle long-term dependencies
- **Transformer (BERT, GPT)** – Modern and powerful

Key NLP Algorithms & Techniques

Text Classification

Used in spam filtering, sentiment analysis.

- **Naive Bayes**
  - Assumes features are independent.
  - **Formula (Bayes' Rule):**
    $P(class|text) = [P(text|class) \times P(class)] / P(text)$
- **SVM**
- **Logistic Regression**
- **BERT-based classifiers**

Named Entity Recognition (NER)

Detects entities like PERSON, LOCATION, ORG.

> "Steve Jobs founded Apple" → ["PERSON", "ORG"]

POS Tagging

Assigns grammatical tags.

> "The dog barked" → [DET, NOUN, VERB]

Dependency Parsing

Identifies syntactic relationships between words.

> "He gave her a book"
> → subject: "He", object: "book", indirect object: "her"

Deep Learning in NLP

RNN (Recurrent Neural Networks)

Handles sequences.

Limitations: Vanishing gradients.

LSTM (Long Short-Term Memory)

Improves RNN by remembering long-term dependencies.

GRU (Gated Recurrent Unit)

Simpler, faster than LSTM with similar performance.

Attention Mechanism

Helps the model focus on important words.

In translation: "I am eating" → "Je mange", focus on "eating" → "mange".

Transformers

Introduced in "Attention Is All You Need" (Vaswani et al., 2017)

- **Architecture:** No recurrence, only attention
- **Used in:** BERT, GPT, T5, RoBERTa, XLNet

Key Transformer Models:

Model	Use
BERT	Bidirectional Encoder Representations (great for understanding)
GPT	Generative Pretrained Transformer (great for text generation)
T5	Text-to-Text Transfer Transformer
DistilBERT	Lightweight BERT
RoBERTa	Robustly Optimized BERT

Semantic Similarity

- Measuring how similar two pieces of text are (used in search engines, chatbots).
- Common metrics: **Cosine Similarity** (on word embeddings).

Formula:

$$\cos(\theta) = (A \cdot B) / (||A|| \ ||B||)$$

Multilingual NLP

- Handling multiple languages in one model (e.g., **mBERT**, **XLM-R**).
- Useful for global applications.

Zero-shot / Few-shot NLP

- NLP without fine-tuning on the target task.
- Example: GPT-4 can perform summarization even without training directly on it.

Prompt Engineering (for LLMs)

- Crafting prompts for better results from generative models like GPT, Claude, etc.
- Very relevant in **Generative AI + NLP**.

Conversational AI Pipelines

- Multi-turn chatbots (e.g., Rasa, Dialogflow)
- Concepts: intent recognition, slot filling, dialogue management

Topic Modeling

- Unsupervised technique to discover hidden topics in a collection of documents.
- Popular algorithms:
    - **LDA (Latent Dirichlet Allocation)**
    - **NMF (Non-negative Matrix Factorization)**

Speech-related NLP

Speech-related NLP focuses on converting spoken language into text and vice versa, enabling natural communication between humans and machines.

ASR (Automatic Speech Recognition)

- **Definition**: Technology that converts human speech into written text.
- **Working**:
    1. **Audio Input** → captured via microphone.
    2. **Preprocessing** → noise reduction, normalization.
    3. **Feature Extraction** → MFCCs, spectrograms.
    4. **Acoustic Model** → maps audio features to phonemes.

5. **Language Model** → predicts the most likely word sequence.

- **Applications**:
  - Voice assistants (Alexa, Siri)
  - Meeting transcription
  - Voice-controlled systems

TTS (Text-to-Speech)

- **Definition**: Technology that converts text into natural-sounding speech.
- **Working**:
  1. **Text Analysis** → splitting into sentences, identifying pronunciation.
  2. **Phonetic Conversion** → converting words to phonemes.
  3. **Prosody Generation** → rhythm, stress, intonation.
  4. **Waveform Generation** → neural vocoders (WaveNet, HiFi-GAN).
- **Applications**:
  - Accessibility tools for visually impaired
  - Audiobook generation
  - Interactive voice bots

Multimodal NLP

Multimodal NLP combines **text** with other modalities such as **images, audio, or video** for richer understanding.

Examples:

- **CLIP** (OpenAI) → Understands images + captions.
- **GPT-4V** → Can process text + images for reasoning.
- **Speech2Text + Image Captioning** → Useful for accessibility.

Applications:

- Content moderation in videos
- Educational tools combining speech + visuals
- Video search with natural language queries

Ethics & Privacy

When processing **audio and text data**, especially in a grammar-scoring or voice-analysis app, ethics and compliance are critical.

GDPR (General Data Protection Regulation) Implications:

- **Data Minimization**: Collect only necessary audio/text data.

- **User Consent**: Must obtain explicit permission before recording.
- **Right to Erasure**: Users can request deletion of their data.
- **Data Security**: Store audio files and transcriptions securely with encryption.
- **Anonymization**: Remove identifiable information from stored text/audio.

Other Considerations:

- **Bias in ASR/TTS**: Ensure models work well across accents and dialects.
- **Misuse Prevention**: Avoid generating harmful or deepfake audio.
- **Transparency**: Inform users how their data will be used.

Evaluation Metrics for NLP

Task	Metric	Formula/Meaning
Classification	Accuracy, Precision, Recall, F1	Measure correctness of predictions
NER, POS, etc.	F1-score	Balances Precision & Recall
Summarization	ROUGE (Recall-Oriented Understudy for Gisting Evaluation)	Compares overlap of n-grams
Translation	BLEU (Bilingual Evaluation Understudy)	Measures match with reference translations

Popular NLP Libraries & Tools

Purpose	Libraries
General NLP	NLTK, spaCy
Deep NLP	Hugging Face Transformers, Flair
Preprocessing	TextBlob, gensim
Word Embeddings	Word2Vec, FastText
OCR	Tesseract, EasyOCR
Datasets	HuggingFace Datasets, Kaggle

NLP Applications

Industry	Use Cases
Healthcare	Clinical report summarization, chatbot assistants
Finance	Document classification, sentiment analysis
E-commerce	Product search, reviews analysis

Industry	Use Cases
EdTech	Essay scoring, AI tutors
Legal	Contract summarization
Social Media	Trend analysis, content moderation

Challenges in NLP

- **Ambiguity**: "I saw her duck" — what does "duck" mean?
- **Sarcasm**: "Oh great, another Monday!" → tone matters
- **Multilingual NLP**
- **Low-resource languages**
- **Bias & fairness in language models**
- **Context understanding**

Future of NLP

- Multilingual Transformers (mBERT, XLM-R)
- Low-resource NLP & Zero-shot learning
- Real-time translation
- Emotion & sentiment-aware assistants
- Ethical and explainable NLP
- Vision + Language models (e.g., GPT-4V)

Summary Checklist

Stage	What to Learn
Basics	Tokenization, stopwords, stemming
Representation	TF-IDF, Word2Vec, BERT
Modeling	Classification, NER, POS
Deep NLP	LSTM, Attention, Transformers
Applications	Chatbots, translation, summarization
Evaluation	F1, BLEU, ROUGE
Tools	spaCy, HuggingFace, NLTK